




# NOORUL HUDA KHAN

 +919990734395  noorulhuda101099@gmail.com  linkedin.com/noorul  https://github.com/THENHKHAN

## Experience

### Fluree — Associate Data Engineer (Full-time)

Aug 2023 – present

Data Engineer

Bangalore, India

- Added functionalities and tools to convert TTL, RDF, and OWL files into organized formats, enhancing overall data management efficiency.
- Enhanced concurrent data processing in a Spark-based application by implementing logically isolated SparkSession instances for each actor using Akka. Applied session isolation techniques, ensuring seamless parallel execution of multiple actors, which improved scalability, performance, and resource efficiency.
- Developed a Python-based tool to convert Microsoft Dynamics 365 Common Data Model (CDM) JSON schemas into JSON-LD format compatible with FlureeDB, enabling semantic data modeling and easier graph database integration.
- Automated batch processing of entity definitions, including schema versioning, RDF class/property generation, and structured ontology linking using CDM standards. Aligned the schema transformation process with FlureeDB's data modeling requirements to support supply chain and inventory data structures in a graph-based context.
- Engineered a robust address normalization pipeline using LibPostal, pandas, and pycountry, introducing Unicode-safe cleaning, duplicate removal, and fallback logic for missing countries, ensuring accuracy and resilience across 1M+ records.
- Optimized runtime by **60%** by caching country lookups with `lru_cache`, replacing slow loops with vectorized operations, and integrating parallel-safe address cleaning via `swifter`.
- I built a PySpark pipeline that uses `SentenceTransformers` to batch-encode millions of records, stores them in a FAISS index, and performs fast nearest-neighbor searches for deduplication. I optimized it with batch processing, distance thresholds, and Spark SQL post-processing to handle **large datasets** in minutes instead of hours.

### Fluree — Data Engineering Intern

Feb 2023 – Aug 2023

Internship

Bangalore, India

- Designed and implemented address parsers in Julia and C, optimizing data extraction and contributing to improved geospatial processing within the system.
- Engineered custom queries, algorithms, and APIs for an internal JSON-LD database using Scala, Spark, and Apache Jena, optimizing data retrieval processes. Engineered custom prompts and trained ChatGPT to understand and execute FlureeDB queries, improving efficiency and user experience.

## Technical Skills

**Languages:** Python, Scala, Spark, Julia, PySpark, Sparql, Java, Javascript

**Backend:** Spring Boot, Fast-API, Familiar with (Flask, Django), SQL, PostgreSQL, MongoDB

**Technologies:** Docker, Databricks, AWS, Azure, Linux, Git, Github, Spring Boot, GraphQL, Postman

**AI Models:** all-MiniLM-L6-v2 (To Generate Embeddings), local LLaMA model (Q & A), mistral-7b-instruct-v0.1.Q4\_K.M Hugging Face model, MCP-Server

## Education

### Aligarh Muslim University

June. 2021 – July 2023

Masters of Computer Applications (MCA) : 8.4 CGPA

AMU, India

### Jamia Millia Islamia

June. 2017 – July 2020

B.sc(Hons) Physics : 8.7 CGPA

JMI, India

## Projects

### Build Real-Time REST APIs with Spring Boot -Blog App | Java, Hibernate, PostgreSQL, Docker

 [Github](#)

- \* Build Rest API's for Blog App Using Spring Boot, Spring Security, JWT , Spring Data JPA (Hibernate), PostgreSQL Database.
- \* Implemented Pagination and Sorting, Login and Sign-up, REST API Validation, Handled exceptions globally as well as specific.

### Task Reminder Application | fastapi, oauth2, react.js, mongodb, javascript

 [Github](#)

- \* Developed a Task Reminder Application with FastAPI, OAuth2, and MongoDB, demonstrating proficiency in backend development.
- \* Implemented a responsive front end using React.js, enhancing user experience and interaction in the Task Reminder Application.